Unraveling with AI Decision: Interpreting and Breaking Down the Black Box using SHAP

Tanishq Varpe
Affiliation
tanishqvarpe@gmail.com

Shravan Raut
Affiliation
shravanraut0202@gmail.com

Rohan Gundgal
Affiliation
gundgalrohan@gmail.com

Abstract: .

Explainable AI (XAI) has the goal of making machine learning models more transparent and reliable by reducing their opacity. Among the techniques proposed to assign feature importance to model predictions are SHapley Additive exPlanations (SHAP) and Local Interpretable Model Agnostic Explanations (LIME), two of the most widely used. This work focuses on the rudiments of the mentioned approaches, with their advantages and disadvantages from the perspective of feature collinearity and model dependence. It offers solutions to real-world problems as well as examples, further increasing the dependability of model explanations for critical decision-making applications in the biomedical case study.

Index terms: interpretability of machine learning, XAI, transparent AI, blackbox models, SHAP, LIME, explainable AI, and trustworthy AI.

INTRODUCTION

In recent years, From financial forecasting and natural language processing to medical diagnosis and tailored healthcare, machine learning (ML) and deep learning approaches have seen tremendous success in recent years. Complex models that frequently outperform conventional statistical techniques, like deep neural networks and ensemble methods, have been a major factor in these advancements. Nevertheless, these models usually function as "black boxes," providing little information about how particular input characteristics affect their predictions. Their adoption is severely hampered by this lack of interpretability, particularly in high-stakes situations where choices have an immediate bearing on safety, human lives, or ethical issues.

A key field of study focused on improving the interpretability and transparency of intricate machine learning models is Explainable Artificial Intelligence (XAI). The objective of XAI is to offer concise, trustworthy, and intelligible explanations that can shed light on the logic underlying model outputs. These justifications are essential for fostering user and AI system trust, meeting ever-tougher regulatory standards, and facilitating model validation and enhancement.

Because of their intuitive interpretation and model-agnostic nature, SHapley Additive exPlanations (SHAP) and Local Interpretable Model-Agnostic Explanations (LIME) have emerged as two of the most popular XAI techniques. In order to fairly distribute "credit" among input features and provide explanations that align with the contributions of features to model predictions, SHAP makes use of cooperative game theory concepts. In contrast, LIME uses interpretable surrogate models, usually linear, to approximate complex model behaviour locally around individual predictions, enabling a case-by-case understanding of feature impact.

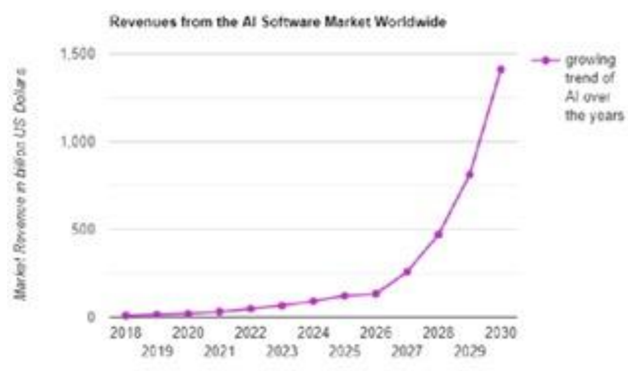

Fig 1: Worldwide AI revenue and growth
Source: A Review of Trustworthy and Explainable Artificial Intelligence (XAI)

Model dependency and feature collinearity are two major practical problems with SHAP and LIME, considering their widespread adoption. Different machine learning algorithms trained on the same data may diverge in explanations regarding feature importance. Similarly, the assumption of independence that grounds both SHAP and LIME can produce unstable and misleading explanations where input features are highly correlated. These issues have the potential to erode user trust and restrict the usefulness of XAI.

In this paper, we provide a detailed review of the methodological underpinnings, main benefits, and inherent drawbacks of SHAP and LIME. We show how feature correlations and model selection impact the outputs of explanations using an empirical case study on biomedical data. We then discuss state-of-the-art techniques such as Modified Index Position (MIP) and Normalized Movement Rate (NMR) that are targeted at assessing and strengthening the robustness and reliability of feature importance rankings in the presence of these challenges. Lastly, we present best practices and recommendations for practitioners and legislators working to field reliable and explainable AI systems in critical domains.

What is XAI?

Explainable AI methods, here termed XAI, are techniques for making the outcomes and actions of machine learning models intelligible to humans. XAI enables appropriate artificial intelligence deployment to sensitive areas by showing which input features drive the model decisions. It engenders users' trust, helps maintain regulatory compliance, and simplifies debugging and model improvement.

Explainable AI, or XAI, is a subset of techniques and methods that provide insights for humans into how such advanced AI and machine learning models make their decisions. In other words, XAI aims at showing the inner logic a model follows to generate its output, while traditional "black-box" models simply output predictions without giving reasons. This transparency will be instrumental in building trust with the user, facilitating regulatory compliance, and providing developers and other stakeholders with the capability to validate, debug, and improve AI models. Particular importance is held by XAI in high-stakes industries like healthcare, finance, and independent systems, wherein it is required that decisions be explainable to ensure their moral and secure use.

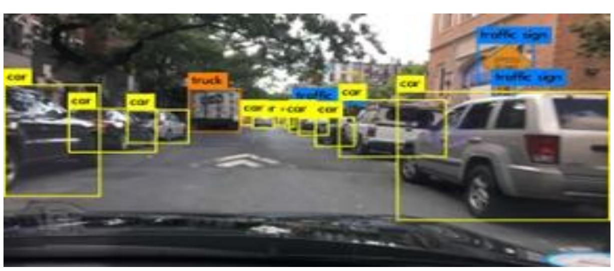

Fig 2: Predicting the object's behaviour: explainable interpretable
Source: A Review of Trustworthy and Explainable Artificial Intelligence (XAI)

XAI methods provide different levels of explanation: local explanations explain why the model made the specific prediction for an instance, while global explanations summarize models' general behaviour. With clear explanations, users are able to understand better how features interact, which input features have the most impacts on model outputs, and the degree of confidence or uncertainty in predictions.

What is SHAP?

Using concepts from cooperative game theory, SHAP calculates each feature's Shapley value, the average marginal contribution of the feature across all subsets of features in the model. This algorithm offers both global explanations across the datasets and local explanations for individual predictions. Its major weakness is that SHAP assumes independence of features, which correlated data would deny, hence distorting attributions. Its other limitation relates to the computational loads, for which approximations have been developed, for instance, Kernel SHAP.

SHAP is an explanation technique under XAI, and it emanates from cooperative game theory with roots in the concept of Shapley values introduced by Lloyd Shapley. The approach views every feature in a machine learning model as a "player" in a cooperative game, with the "payout" being the model's predicted output.

Key Concept: SHAP is calculated by averaging the marginal contributions of each feature across all possible coalitions of features. This methodology ensures that the features will have an equitable share of the payout with respect to their influence. SHAP explains both local- per instance-and global- overall model behaviours by aggregating these contributions.

METHODOLOGY

In the research, several machine learning models are compared, such as decision trees, logistic regression, gradient boosting, and support vector machines, which will be trained on biomedical data from the UK Biobank to classify the occurrence of myocardial infarction. Feature importance will be interpreted using SHAP and LIME. Explanation stability regarding model choice and feature collinearity will be assessed by two metrics, NMR and MIP, for evaluating and correcting feature importance rankings. Making sense of model predictions using SHAP proceeds according to a workflow that is depicted in Figure 1. An appropriate dataset can be assembled and pre-processed and then becomes the basis on which the model is trained. Once the dataset has been assembled and pre-processed, selecting an appropriate machine learning model-one that learns the patterns and relationships represented in the data-occurs.

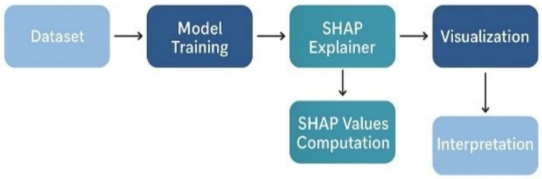

Fig 3: Methodology Flow chart

Following the fitting of the model, a SHAP Explainer is used to calculate the Shapley value of the features. The intent of the computation is to understand how each input variable contributes positively or negatively to the model's output for each sample. These SHAP values are then visualized using several specialized plots: bar graphs, summary plots, and force plots. These visualizations provide both a global view of the estimated feature importance across the entire dataset and a local view that estimates the impact of individual features on a specific prediction. The interpreted visual outputs are then elicited for actionable interpretations, such as questions regarding the major drivers of model decisions to validate model behaviour, or to inform further the process of model refinement. Using this approach, researchers can generate transparent, interpretable, and meaningful descriptions for complex machine learning models supportive of trust and responsible use.

SYSTEM ARCHITECTURE

The proposed system architecture is a layered design that will support the complete workflow of data processing, model training, evaluation, and explainability using SHAP and LIME.

1. Data Acquisition & Storage Layer: Supply and store structured datasets like CSV, Excel, SQL, biomedical, sensor; perform validation, filtering, and schema checking; ensure machine learning pipeline compatibility.

2. Data Preprocessing Layer: Prepares datasets for modeling by handling missing values, outliers, noise reduction, encoding categorical variables, feature engineering, normalization, and splitting data into training, testing, and validation sets.

3. Machine Learning Model Layer: Train predictive models, which could include Decision Trees, Logistic Regression, SVM, Gradient Boosting, and Neural Networks using preprocessed data. The goal is to store the model artifacts and generate predictions with probability scores, allowing for explainability.

4. XAI Layer: Integrating SHAP and LIME for both global and local explanations. SHAP calculates feature importance and summarizes it with summary plots, beeswarm, bar, and force plots. LIME builds local surrogate models and does feature sensitivity analysis.

5. Visualization & Reporting Layer: This will convert the explanations into easily understandable visual and textual forms, like dashboards, comparative charts, heatmaps, scatter plots, and exportable reports (PDF/HTML), while providing optional interactive interfaces to formulate queries.

6. User Interaction Layer: This layer allows users to upload data and view predictions apart from comparing SHAP and LIME explanations. It targets doctors, business users, analysts, and domain experts.

CASE STUDY

The dataset included 1500 subjects with 10 predictive features. While the models returned similar accuracy, they returned features ranked with different levels of importance. The two key predictors identified were cholesterol and body mass index, which are highly collinear with one another and thus affect explanations returned from SHAP and LIME. Application of NMR analysis showed that the gradient boosting models returned the most stable rankings of feature importance. The incorporation of MIP adjustments contributed to terminological consistency in terms of interpretation found in MIP models, due to consideration of multicollinearity effects and implications for trusting that level of feature selection ranked higher than other features.

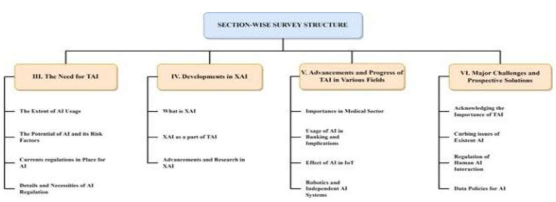

Fig 4: Section wise survey structure
Source: A Review of Trustworthy and Explainable Artificial Intelligence (XAI)

Practical Case Study

A bank uses SHAP in understanding a predictive model for customer churn. Some of the main takeaways include:
Service Usage, Custoxmer Age were ranked as the most important features in customer churn by NSC.
SHAP plots show how both features have non-linear effects, and interactions with one another, that might not be

captured through the traditional longitudinal feature importance dashboard.

• With an explanation for each individual on prediction models, customer service teams will proactively focus on targeting and engaging customers most at-risk of churning to increase retention rates.

The case illustrates how SHAP amplifies the twin benefits of explainability and actionable insight, transforming an intangible AI algorithm into AI-powered decision-making with a proper consideration of business objectives.

II. Recent Progress and Issues Associated with XAI

The models within machine learning, mainly deep neural networks and ensemble models, have managed to achieve remarkable accuracy in decision-making contexts. Yet, the lack of interpretability of this method is a serious problem in technical fields, like healthcare, banking, and environmental science.

Thus, XAI is designed to bring transparency/trust in the decision-making of models through ethical and regulatory demands for explainability. Current methods of XAI can easily be reduced to two groups: that of either transparent model development, which embeds interpretability natively -for example, models based on rules, decision trees, etc.-, and that of post-hoc explanation methods, explaining the predictions of complex black-box models after they have been developed.

III. Strengths and Weaknesses of SHAP and LIME

SHAP provides both global and local explanations by attributing the value of a variable in the predictive model using Shapley values derived from cooperative game theory. Flexibility is also one of its strengths, since it applies to any model and data structure. Moreover, SHAP has summary plots and force plots to visually communicate the contribution of each feature to an output of the model to either the influenced variable or the prediction.

• LIME generates straightforward and interpretable surrogate models around a specific prediction for fast local explanation through visual interpretation. However, the major limitations of LIME involve assumptions such as feature independence and local linearity, which have a potential to be unreliable in the presence of non-linearity or correlated features.

DISCUSSION

Most traditional methods of explaining machine learning algorithms rely on global variable importance plots, which can only provide global and static assessments of the importance of features under that model. In contrast, SHAP provides a more granular and interpretable understanding of how each predictor is contributing to not only the overall model but also individual predictions, somewhat like regression coefficients but at a finer-grained level and in context. In this way, it becomes possible for practitioners to examine both global trends and localized effects of features, knowledge that does not typically come from traditional statistical models.

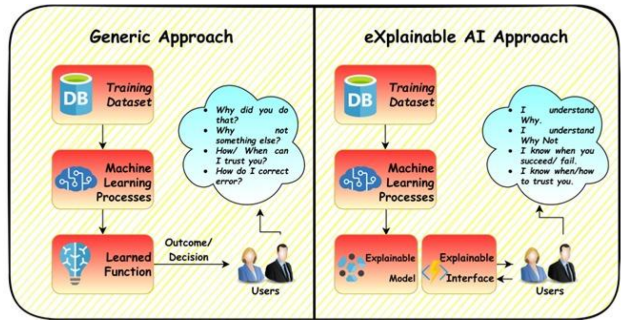

Fig 5: Generic vs XAI approach
Source: A Review of Trustworthy and Explainable Artificial Intelligence (XAI)

Some of the main advantages of SHAP are that these contributions can be visualized and positioned spatially, as many of our earlier studies on hazard prediction and biomedical analysis have shown. These SHAP maps will show that the importance of a feature, such as NDVI or drainage density, is constantly changing over geography or populations, thus enabling better risk assessments and more precise interventions. Furthermore, if we couple this with an interactive visualization tool-a web-based GIS, for example-stakeholders can query and project model outputs for specific regions or instances directly without having to wait for research results to inform any decisions.

Constraining our exploration of SHAP to its benefits, SHAP approaches cadre some limitations. This is because AI-based approaches are computationally intensive and can still be influenced by concerns such as collinearity among predictors and the absence of temporal information in the input data. Mapping susceptibility and exposure is a big leap forward, but risk characterization and full risk assessment would require adding time-dependent and intensity components as well. Work was also supported, in part, by even grants and programs aimed at promoting transparency and the ethical application of AI in sweet-like high stakes areas. The support received from different interdisciplinary teams ranging from data science to machine learning.

RESULTS

1. Scope of Model Interpretability and Explanation:
SHAP provides the capability for both global and local interpretability: to know the overall feature importance, but also how much each feature alone contributed to the prediction on an individual prediction level.

LIME is primarily a local interpretability technique that explains individual predictions by training simple surrogate models around the vicinity of any given instance.

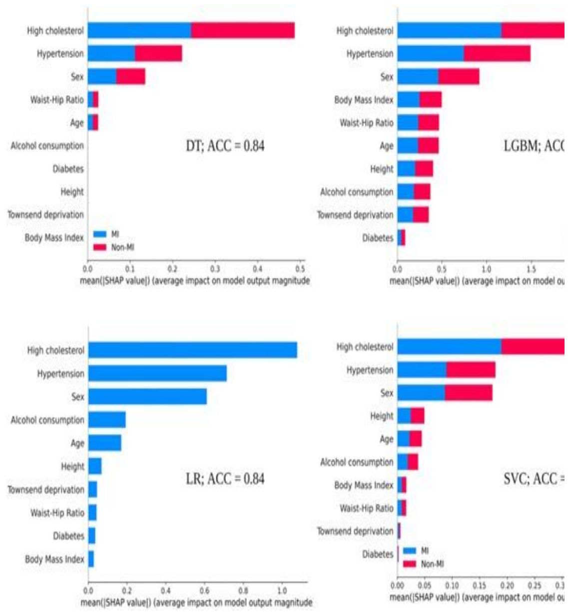

Fig 6: SHAP output to explain the four models globally. DT: decision tree; LGBM: light gradient-boosting machine; LR: logistic regression; SVC: support vector machines classifier; ACC: accuracy; MI: myocardial infarction.
Source: A Perspective on Explainable Artificial Intelligence Methods: SHAP and LIME.

2. Implementation and Applicability:
SHAP and LIME are post-hoc, model-agnostic explainer methods utilized after training the model, and they can be applied to any machine learning model.
SHAP does particularly well with more complex models, like deep neural networks and ensemble methods. On the other hand, LIME is lighter computationally and allows for either studying simpler models or obtaining some quick, instancelevel explanations.

3. Explanation Type and Visualization:
SHAP generates numerical and visual explanations, including summary, waterfall, and bees warm plots, which can be useful for understanding the features' contribution in both magnitude and direction.
LIME does textual and visual explanations, usually as bar plots, but for an individual prediction, not giving a general global picture.

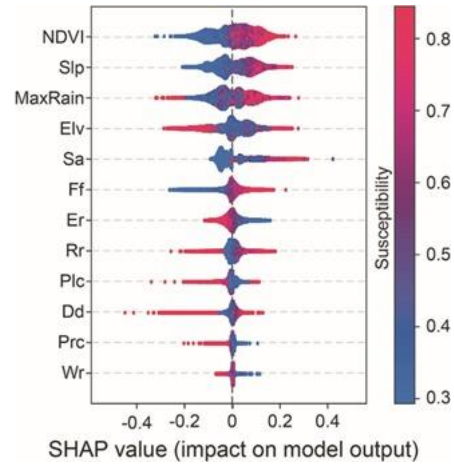

Fig 7: The SHAP value distribution for each variable against the susceptibility. Each dot corresponding to a specific catchment, the colour map showed the final susceptibility.
Source: On the use of explainable AI for susceptibility modeming: Examining the spatial pattern of SHAP values

4. Stability, Consistency, and Limitations:
SHAP produces more stable and consistent explanations, important for sensitive applications requiring consistency and reliability.
Instabilities in LIME may arise due to random sampling, reliance on local surrogate fitting, and variability across multiple runs.
Both have challenges with collinearity or high-dimensional data, although this is less pronounced for SHAP especially in the presence of non-linear feature interactions.

5. Performance in Real-World Studies:
Various clinical and biomedical studies have shown high predictive performance using SHAP and LIME on top of machine learning models: for instance, 97.6% accuracy, 95.8% precision, 97% recall, and 96.8% F1 score in disease diagnosis tasks.
SHAP is usually preferred as it provides a quantitative ranking of features and a more global view of model behaviour.
On the other side, LIME is appreciated for its flexibility in use when interpreting single predictions.

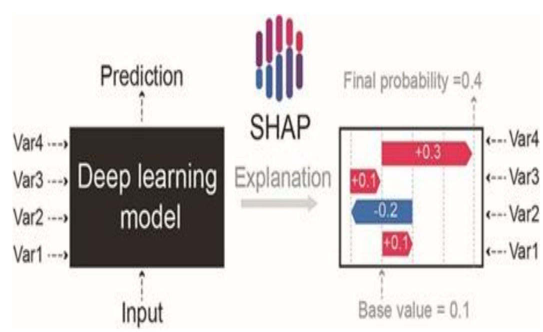

Fig 8: An illustration demonstrating the SHAP-explained deep learning models

Source: On the use of explainable AI for susceptibility modelling: Examining the spatial pattern of SHAP values

6. User Recommendations:

SHAP is recommended in cases when it is necessary to obtain broad, reliable, and consistent interpretability of very complex models.

LIME is more practical when quick, instance-level insights are desired and/or computational resources are limited.

Both techniques can be used to establish trust in and provide better insight into explainable models, normally applied in high-stake domains such as medical diagnosis, financial forecasting, and environmental risk estimation.

CONCLUSION

While SHAP and LIME are very useful in XAI, there are certain limitations that should be taken seriously. Model dependence and feature correlation call for better a posteriori adjustment and the employment of various models in order to ensure the reliability of the interpretations. Further research will aim at the development of hybrid methods sensitive to data dependencies that will foster trustworthy and practical explainability of important applications.

ACKNOWLEDGEMENT

This paper would not be possible without the various supports and collaborative efforts of individuals and organizations. We would like to thank the scientific community who have done extensive work on explainable AI. Their works were foundational to undertaking this study. We also want to extend my gratitude to the developers and maintainers of the open-access libraries like SHAP and LIME. Their tools and documentation have been essential to the analyses and visualizations presented here.

We would like to extend particular thanks to our academic mentors, colleagues, and collaborators for their valuable feedback and discussions while this paper was developed and written. Without contributions and shared resources from experts in healthcare, environmental science, and geospatial technology, the datasets and case studies here presented would not have been possible. The institutions and research centers that gave access to computational resources and data repositories are appreciated by the authors.

This work was also supported by grants and initiatives with a view to encouraging transparency and ethics in the use of AI in critical sectors. The support from diverse teams in data science, machine learning, risk analysis, and public policy has considerably enriched the perspectives included in this research.

Finally, sincere thanks go to the reviewers and editors, whose comments and suggestions increased the quality and clarity of the manuscript. The dedication of the growing XAI community to open science and continuous improvement is an inspiration for and guide towards responsible and impactful AI research.

REFRENCES

1. A Perspective on Explainable Artificial Intelligence Methods: SHAP and LIME
Ahmed M. Salih, Zahra Raisi-Estabragh, Ilaria Boscolo Galazzo, Petia Radeva, Steffen E. Petersen, Karim Lekadir, and Gloria Menegaz,
Advanced Intelligent Systems, 2025.
Available: https://arxiv.org/html/2305.02012v3
2. Interpreting artificial intelligence models: a systematic review of LIME and SHAP:
V. Vimbi et al., PLOS Computational Biology, 2024.
Available: https://pmc.ncbi.nlm.nih.gov/articles/PMC10997568/
3. A Comparative Analysis of SHAP and LIME on Various ML Models:
M. M. Hasan et al., Journal of Science and Technology Research, 2023.
Available: https://jstr.bousst.edu.bd/index.php/jstr/article/view/5/7
4. A Comparative Analysis of Explainable AI Techniques for Model Interpretability :
Y. Swathi et al., 2023 International Conference on Artificial Intelligence, IEEE.
Available: https://ieeexplore.ieee.org/document/10266190/
5. Unlocking Machine Learning Model Decisions:
Journal of Electrical Systems, 2024.
Available: https://journal.esrgroups.org/jes/article/view/1768